Log Concave Polynomials II: High-Dimensional Walks and an FPRAS for Counting Bases of a Matroid

Nima Anari, Kuikui Liu, Shayan Oveis Gharan, Cynthia Vinzant

Presented by Alex Albors Juez and Rohan Mukherjee

CSE 521, University of Washington

December 14th, 2023

Table of contents

Matroids

A *matroid* M is a pair $M = (X, \mathcal{I})$ where X is a finite set and $\mathcal{I} \subseteq 2^X$ so that the following holds:

(i) *Non-emptyness*: $\emptyset \in \mathcal{I}$

(ii) *Monotonicity*: If $Y \in \mathcal{I}$ and $Z \subseteq Y$ then $Z \in \mathcal{I}$.

(iii) *Exchange property*: If $Y, Z \in \mathcal{I}$ and $|Y| < |Z|$, then for some $x \in Z \backslash Y$ we have $Y \cup \{x\} \in \mathcal{I}$

Definition (basis)

Let $M = (X, \mathcal{I})$ be a matroid. A maximal independent set $B \in \mathcal{I}$ is called a *basis* of X. All basis elements have the same size, and their size is called the *rank* of the matroid.

Example: The Acyclic subsets of a graph (forests) form a matroid, called a *graphic matroid*.

Bases exchange walk

Procedure:

1. Start with a basis element B.
2. Drop a random element i from B. Pick j uniformly at random from $\{1, \ldots, n\}$, and try adding it to $B \backslash \{i\}$. Do it until we can.
3. Repeat step 2.

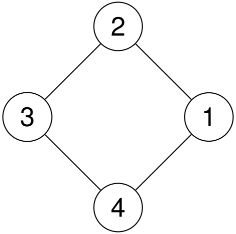

Figure 1: Graph C_4 corresponding to a rank 3 graphic matroid

History

- 30 years ago, Mihail and Vazirani conjectured that the bases exchange walk mixes in polynomial time.
- Polynomial mixing time corresponds to being able to count bases in polynomial time (Approximate sampling and approximate counting are equivalent in this scenario [2, JVV86]).
- Barvinok and Samorodnitsky designed a randomized algorithm that gives a $\log(n)^r$ approx. factor for a matroid with n elements and rank r [4, BS07].
- In Log-concave polynoimals I, Gharan et al. give a deterministic algorithm that returns an e^r approximation factor.[3, AKOV18]
- In this paper, Gharan et al. give a randomized algorithm yielding a $1 \pm \epsilon$ approximation factor in polynomial time.

Main theorem

Theorem (1.1)

Let $\mu : 2^{[n]} \to \mathbb{R}_{\geq 0}$ *be a* d*-homogeneous strongly log concave probability distribution. If* P_μ *denotes the transition probability matrix of* M_μ *and* $X(k)$ *denotes the set of size-*k *subsets of* $[n]$ *which are contained in some element of* $\mathrm{supp}(\mu)$*, then for every* $0 \leq k \leq d-1$*,* P_μ *has at most* $|X(k)| \leq \binom{n}{k}$ *eigenvalues of value* $> 1 - \frac{k+1}{d}$*. In particular,* M_μ *has spectral gap at least* $1/d$*, and if* $\tau \in \mathrm{supp}(\mu)$ *and* $0 < \epsilon < 1$*, the total variation mixing time of the Markov chain* M_μ *started at* τ *is at most* $t_\tau(\epsilon) \leq d \log(\frac{1}{\epsilon \mu(\tau)})$*.*

Simplicial Complexes

Definition

- A set $X \subseteq 2^{[n]}$ is called a simplicial complex if whenever $\sigma \in X$ and $\tau \subset \sigma$, we have $\tau \in X$.
- Elements of X are called faces, and the dimension of a face $\tau \in X$ is defined as $\dim(\tau) = |\tau|$.
- A face of dimension 1 is called a *vertex*, and a face of dimension 2 is called an *edge*.
- Define $X(k) = \{ \tau \in X \mid \dim(\tau) = k \}$ to be the collection of degree-k faces of X.

Examples

Any (undirected) graph $G = (V, E)$ is an example of a simplicial complex.

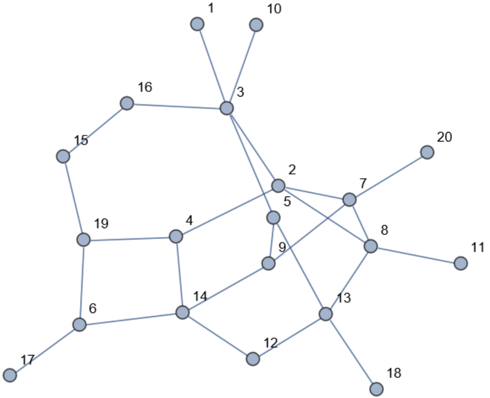

Figure 2: Example of a simplicial complex

Examples (contd.)

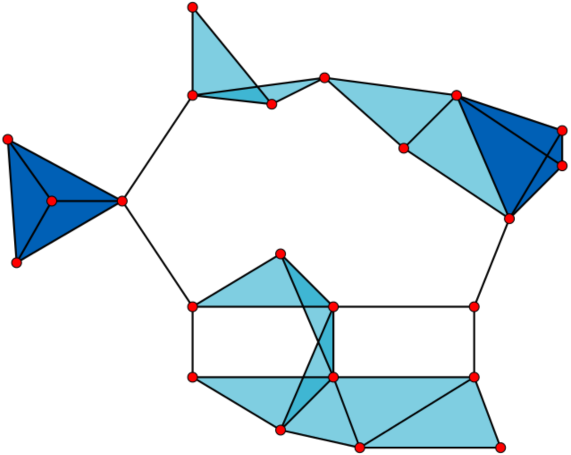

Figure 3: Example of a simplicial complex

Definitions contd.

- A simplicial complex X is pure if all maximal (w.r.t. inclusion) faces have the same dimension.
- The link of a face $\tau \in X$ is defined by $X_\tau = \{\sigma \setminus \tau \mid \sigma \in X,\ \tau \subset \sigma\}$. Importantly, if X is pure of dimension d and $\tau \in X(k)$, then X_τ is pure of dimension $d - k$.
- Can equip a weight function $w : X \to \mathbb{R}_{>0}$ to X by assigning a positive weight to each face of X. Say a weight function $w : X \to \mathbb{R}_{>0}$ is balanced if for any $\tau \in X$,
$$w(\tau) = \sum_{\substack{\sigma \in X(k+1) \\ \tau \subset \sigma}} w(\sigma)$$
- Notice that we can equip X with a (balanced) weight function by assigning its maximal faces weights and then assigning weights to the rest of the faces inductively.

Weights contd.

- Any (balanced) weight function on X induces a weighted graph on the vertices of X as follows: the 1-skeleton of X is the (weighted) graph $G = (X(1), X(2), w)$ where w has been restricted from X to $X(2)$. In this case $w(v)$ for $v \in X(1)$ is the weighted degree of v.

d-homogeneous polynomials

A polynomial $p \in \mathbb{R}[x_1, \ldots, x_n]$ is d-homogeneous if $p(\lambda x_1, \ldots, \lambda x_n) = \lambda^d p(x_1, \ldots, x_n)$ for every $\lambda \in \mathbb{R}$. Notice in this case that,

$$\sum_{k=1}^{n} x_k \partial_k p(x) = d \cdot p(x)$$

Example. Consider $p(x, y, z) = xyz^2 + x^2yz$. Then,

$$\begin{aligned}\sum_{k=1}^{3} x_k \partial_k p(x) &= (xyz^2 + 2x^2yz) + (xyz^2 + x^2yz) + (2xyz^2 + x^2yz) \\ &= 4xyz^2 + 4x^2yz\end{aligned}$$

Constructing Simplicial Complexes from Polynomials

From a d-homogeneous $p \in \mathbb{R}_{\geq 0}[x_1, \ldots, x_n]$ $p(x) = \sum_S c_S x^S$, can construct a (weighted) simplicial complex X^p by doing the following: include a d-dimensional face S with weight $w(S) = c_S$ and include all subsets of these maximal faces inductively.

Visuals

This polynomial yields the above (weighted) simplex where each tetrahedral face has weight 1:

$$p(x_1, \ldots, x_7) = x_1x_2x_3x_4 + x_3x_5x_6x_7$$

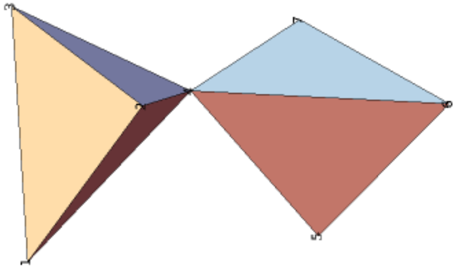

Figure 4: Two Tetrahedrons Glued Together

Roadmap

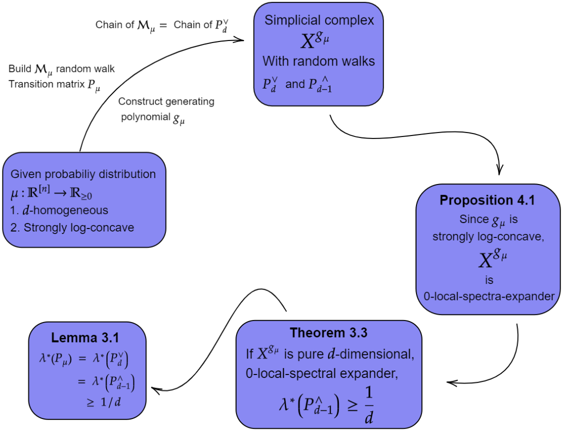

Figure 5: Roadmap

Log-concave polynomial identities

Definition

A polynomial $p \in \mathbb{R}_{\geq 0}[x_1, \ldots, x_n]$ is log-concave if $\log p$ is concave, equivalently if

$$\nabla^2 \log p = \frac{p \cdot (\nabla p)^2 - (\nabla p)(\nabla p)^T}{p^2}$$

is NSD. For convience, define $p(x) \equiv 0$ to be log-concave.

Log-concave properties contd.

- By Cauchy's interlacing theorem, if p is log-concave then $p \cdot (\nabla^2 p)$ has at most one positive eigenvalue at any $x \in \mathbb{R}^n_{>0}$.
- Since p has nonnegative coefficients, log-concavity is equivalent to $\nabla^2 p \preceq \frac{(\nabla p)(\nabla p)^T}{p}$, so in this case $\nabla^2 p$ has at most 1 one positive eigenvalue.
- Turns out converse is true too: if p is a degree d-homogeneous polynomial in $\mathbb{R}[x_1, \ldots, x_n]$, and $(\nabla^2 p)(x)$ has at most one positive eigenvalue for all $x \in \mathbb{R}^n_{>0}$, then p is log-concave.

Definition

A polynomial $p \in \mathbb{R}[x_1, \ldots, x_n]$ is strongly log concave if for all $k \geq 0$ and all $\alpha \in [n]^k$, we have $\partial^\alpha p$ is log-concave (i.e., all sequences of partials are log-concave).

Markov Chains and Random Walks

- A Markov Chain is a triple (Ω, P, π) where Ω denotes a finite state space, $P \in \mathbb{R}_{\geq 0}^{\Omega \times \Omega}$ is a transition probability matrix. That is,

$$P(i,j) = P_{ij} = \mathbb{P}(X_{t+1} = j \mid X_t = i)$$

  It follows that the matrix is stochastic, such that $P\mathbf{1} = \mathbf{1}$. Finally, $\pi \in \mathbb{R}_{\geq 0}^{\Omega}$ denotes the stationary distribution of the chain ($\pi P = \pi$).
- The Markov Chain (Ω, P, π) is reversible if

$$\pi(\tau)P(\tau,\sigma) = \pi(\sigma)P(\sigma,\tau)$$

  for all $\tau, \sigma \in \Omega$.

Markov Chains and Random Walks continued

- For any reversible Markov chain (Ω, P, π), the largest eigenvalue of P is 1 (Perron-Fröbenius Theorem). We let $\lambda^*(P) = \max\{|\lambda_2|, |\lambda_n|\}$. The *spectral gap* of the Markov chain is $1 - \lambda^*(P)$.

Theorem (2.9, (DS))

For any reversible irreducible Markov chain (Ω, P, π), $\epsilon > 0$, *and any starting state*

$$t_\tau(\epsilon) \leq \frac{1}{1 - \lambda^*(P)} \cdot \log\left(\frac{1}{\epsilon\pi(\tau)}\right)$$

where

$$t_\tau(\epsilon) = \min\left\{\, t \in \mathbb{N} \mid \left\|P^t(\tau, \cdot) - \pi\right\|_1 \leq \epsilon \,\right\}$$

Setting the stage

- Consider a pure d-dimensional complex X with a balanced weight function $w : X \to \mathbb{R}_{>0}$.
- Going to define a bipartite graph G_k with one side $X(k)$ and the other side $X(k+1)$. Connect $\tau \in X(k)$ to $\sigma \in X(k+1)$ with an edge of weight $w(\sigma)$ iff $\tau \subset \sigma$. Consider simple random walk on G_k: choose a neighbor proportional to the weight of the edge connecting the two vertices.

Examples

- One on $X(k)$ called $P_k^{\wedge}$, where given $\tau \in X(k)$ you take two steps of the walk in G_k to transition to the next k-face w.r.t. the $P_k^{\wedge}$ matrix.
- One on $X(k+1)$ called $P_{k+1}^{\vee}$, where given $\sigma \in X(k+1)$ you take two steps to transition to the next $k+1$ face w.r.t. $P_{k+1}^{\vee}$.

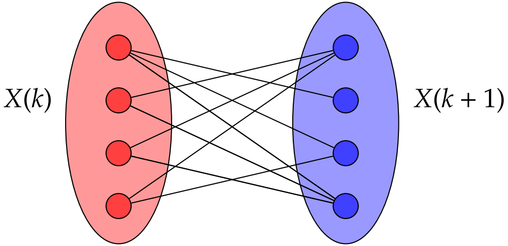

Figure 6: Bipartite graph G_k

Values of the transition matrix

$$P_k^{\wedge}(\tau, \tau') = \begin{cases} \frac{1}{k+1} & \text{if } \tau = \tau' \\ \frac{w(\tau \cup \tau')}{(k+1)w(\tau)} & \text{if } \tau \cup \tau' = X(k+1) \\ 0, & \text{otherwise} \end{cases}$$

$$P_{k+1}^{\vee}(\sigma, \sigma') = \begin{cases} \sum_{\tau \in X(k);\ \tau \subset \sigma} \frac{w(\sigma)}{(k+1)w(\tau)} & \text{if } \sigma = \sigma' \\ \frac{w(\sigma')}{(k+1)w(\sigma \cap \sigma')} & \text{if } \sigma \cap \sigma' = X(k) \\ 0, & \text{otherwise} \end{cases}$$

Note that both random walks are reversible with the same stationary distribution:

$$w(\tau)P_k^{\wedge}(\tau, \tau') = w(\tau')P_k^{\wedge}(\tau', \tau) \quad \text{and} \quad w(\sigma)P_{k+1}^{\vee}(\sigma, \sigma') = w(\sigma')P_{k+1}^{\vee}(\sigma', \sigma)$$

Proving $\lambda^*(P_d^\wedge) = \lambda^*(P_{d-1}^\vee)$

Fact (Useful)

Let $A \in \mathbb{R}^{n\times k}$ and $B \in \mathbb{R}^{k\times n}$ be arbitrary matrices. Then, non-zero eigenvalues of AB are equal to non-zero eigenvalues of BA with the same multiplicity.

Lemma (3.1)

For any $1 \le k \le d$, $P_k^\wedge$ and $P_{k+1}^\vee$ are stochastic, self-adjoint w.r.t. the ω-induced inner product, PSD, and have the same (with multiplicity) non-zero eigenvalues.

Proving $\lambda^*(P_d^\wedge) = \lambda^*(P_{d-1}^\vee)$

Proof.

Since G_k is bipartite, we may write the transition of the random walk on G_k as

$$P_k = \begin{bmatrix} 0 & P_k^\downarrow \\ P_k^\uparrow & 0 \end{bmatrix}$$

Note that $P_k^\uparrow$ and $P_k^\downarrow$ are stochastic matrices. Then we see that

$$P_k^2 = \begin{bmatrix} P_k^\downarrow P_k^\uparrow & \\ & P_k^\uparrow P_k^\downarrow \end{bmatrix}$$

It is easy to see P_k^2 is PSD and stochastic. But now we note that $P_k^\wedge = P_k^\downarrow P_k^\uparrow$ and $P_{k+1}^\vee = P_k^\uparrow P_k^\downarrow$ and we're done. $\square$

Looking at $P_1^\wedge$

- $P_1^\wedge$ is the transition probability matrix of the simple $(1/2)$-lazy random walk on the weighted 1-skeleton of X where the weight of each edge $e \in X(2)$ is $w(e)$.
- Also consider the non-lazy variant of this random walk, given by the transition matrix $\widetilde{P}_1^\wedge = 2(P_1^\wedge - I/2)$
- Similarly, for any face $\tau \in X(k)$, we define the upper random walk on the faces of the link X_τ. Specifically, let $P_{\tau,1}^\wedge$ denote the transition matrix of the upper walk, as above, on the 1-dimensional faces of X_τ, and $\widetilde{P}_{\tau,1}^\wedge = 2(P_{\tau,1}^\wedge - I/2)$ be the transition matrix for the non-lazy version.

Definition (Local Spectral Expanders, KO18)

For $\lambda > 0$, a pure d-dimensional weighted complex (X, w) is a λ-local-spectral-expander if for every $0 \le k < d - 1$, and for every $\tau \in X(k)$, we have $\lambda_2(\widetilde{P}_{\tau,1}^\wedge) \le \lambda$.

Theorem 3.3

Theorem

Let (X, w) be a pure d-dimensional weighted 0-local spectral expander and let $0 \leq k < d$. Then for all $-1 \leq i \leq k$, $P_k^\wedge$ has at most $|X(i)| \leq \binom{n}{i}$ eigenvalues of value $> 1 - \frac{i+1}{k+1}$, where (by convention) $X(-1) = \emptyset$ and $\binom{n}{-1} = 0$. In particular, the second largest eigenvalue of $P_k^\wedge$ is at most $\frac{k}{k+1}$.

Lemma

$P_k^\wedge \preceq \frac{k}{k+1} P_k^\vee + \frac{1}{k+1} I$ for all $0 \leq k < d$.

From log-concavity to Local Spectral Expanders

Theorem (Proposition 4.1)

Let $p \in \mathbb{R}[x_1, \ldots, x_n]$ be a multiaffine homogeneous polynomial with nonnegative coefficients. If p is strongly log-concave, then (X^p, w) is a 0-local-spectral-expander, where $w(S) = c_S$ for every maximal face $S \in X^p$

- Let $p_\tau = \left(\prod_{i \in \tau} \partial_i\right) p$

Lemma (4.2)

For any $0 \le k \le d$, and any simplex $\tau \in X^p(k)$, $w(\tau) = (d-k)!p_\tau(\mathbf{1})$.

From log-concavity to Local Spectral Expanders

Lemma (Lemma 4.2)

For any $0 \leq k \leq d$*, and any simplex* $\tau \in X^p(k)$, $w(\tau) = (d-k)! p_\tau(\mathbf{1})$.

Proof of Lemma.

Induction on $d-k$. If $\dim(\tau) = d$, then $p_\tau = c_\tau$, and done. So suppose statement holds for $\sigma \in X^p(k+1)$ and fix simplex $\tau \in X^p(k)$. Then,

$$w(\tau) = \sum_{\substack{\sigma \in X^p(k+1) \\ \tau \subset \sigma}} w(\sigma) = \sum_{i \in X^p_\tau(1)} w(\tau \cup \{ i \})$$

Since $\partial_i p_\tau = 0$ for $i \notin X^p_\tau(1)$, we have

$$w(\tau) = (d-k-1)! \sum_{i \in X^p_\tau(1)} p_{\tau \cup \{ i \}}(\mathbf{1}) = (d-k-1)! \sum_{i=1}^{n} \partial_i p_\tau(\mathbf{1}) = (d-k)! p_\tau(\mathbf{1})$$

Where the last equality holds by Euler's identity. □

Proof of Proposition 4.1

Since p is strongly log-concave, $\nabla^2 p(\mathbf{1})$ has at most one positive eigenvalue. Let

$$\tilde{\nabla}^2 p = \frac{1}{d-k-1}(\mathrm{diag}(\nabla p))^{-1}\nabla^2 p(\mathbf{1})$$

Claim: $\tilde{\nabla}^2 p = \tilde{P}^{\wedge}_{\tau,1}$. Note that

$$\tilde{P}^{\wedge}_{\tau,1}(i,j) = \frac{w_\tau(\{\, i,j \,\})}{w_\tau(\{\, i \,\})} = \frac{w(\tau \cup \{\, i,j \,\})}{w(\tau \cup \{\, i \,\})}$$

While,

$$(\tilde{\nabla}^2 p)(i,j) = \frac{(\partial_i \partial_j p)(\mathbf{1})}{(d-k-1)(\partial_i p)(\mathbf{1})}$$

By lemma, equal.

Proof contd.

Since p has nonnegative coefficients, the vector $(\nabla p)(\mathbf{1})$ has nonnegative entries which implies $\mathrm{diag}(\nabla p)(\mathbf{1}) \geq 0$. Fact: if $B \geq 0$ and A has (at most) k positive eigenvalues then BA has at most k positive eigenvalues. Since $(\nabla^2 p)(\mathbf{1})$ has at most 1 positive eigenvalue, $\tilde{\nabla}^2 p$ has at most 1 positive eigenvalue by the fact. Thus, $\tilde{\nabla}^2 p = \tilde{P}^{\wedge}_{\tau,1}$ has at most one positive eigenvalue, so $\lambda_2(\tilde{P}^{\wedge}_{\tau,1}) \leq 0$. $\square$

Generating polynomial of μ and $\mathcal{M}_\mu$

- Let $\mu : 2^{[n]} \to \mathbb{R}_{\geq 0}$ be a probability distribution. Assing a multiaffine polynomial with variables $x_1 \dots, x_n$ to μ:

$$g_\mu(x) = \sum_{S \subset [n]} \mu(S) \cdot \prod_{i \in S} x_i$$

- Say μ is d-homogeneous if g_μ is d-homogeneous, and (strongly) log-concave if g_μ is.
- We can define a random walk $\mathcal{M}_\mu$ by the following: We take the state space of M_μ to be the support of μ, namely $\mathrm{supp}(\mu) = \{S \subseteq [n] \,|\, \mu(S) \neq 0\}$. For $\tau \in \mathrm{supp}(\mu)$, first we drop an element $i \in \tau$, chosen uniformly at random from τ. Then, among all sets $\sigma \supseteq \tau \setminus \{i\}$ in the support of μ, we choose one with probability proportional to $\mu(\sigma)$.

Proof of Theorem 1.1

Theorem (1.1)

Let $\mu : 2^{[n]} \to \mathbb{R}_{\geq 0}$ be a d-homogeneous strongly log concave probability distribution. If P_μ denotes the transition probability matrix of M_μ and $X(k)$ denotes the set of size-k subsets of $[n]$ which are contained in some element of $\mathrm{supp}(\mu)$, then for every $0 \leq k \leq d-1$, P_μ has at most $|X(k)| \leq \binom{n}{k}$ eigenvalues of value $> 1 - \frac{k+1}{d}$. In particular, M_μ has spectral gap at least $1/d$, and if $\tau \in \mathrm{supp}(\mu)$ and $0 < \epsilon < 1$, the total variation mixing time of the Markov chain M_μ started at τ is at most $t_\tau(\epsilon) \leq d \log(\frac{1}{\epsilon \mu(t)})$.

Proof.

Let μ be a d-homogeneous strongly log-concave distribution, and let P_μ be the transition probability matrix of the chain M_μ. By Theorem 2.9, it is enough to show that $\lambda^*(P_\mu) \leq 1 - \frac{1}{d}$. Observe that the chain M_μ is exactly the same as the chain $P_d^\vee$ for the simplicial complex X^{g_μ} defined above. Therefore, $\lambda^*(P_\mu) = \lambda^*(P_d^\vee) = \lambda^*(P_{d-1}^\wedge)$, where the last equality follows by Lemma 3.1. Since g_μ is strongly log-concave, by Proposition 4.1, X^{g_μ} is a 0-local-spectral-expander. Therefore, by Theorem 3.3, $\lambda^*(P_{d-1}^\wedge) \leq 1 - \frac{1}{(d-1)+1} = 1 - \frac{1}{d}$. $\square$

Roadmap

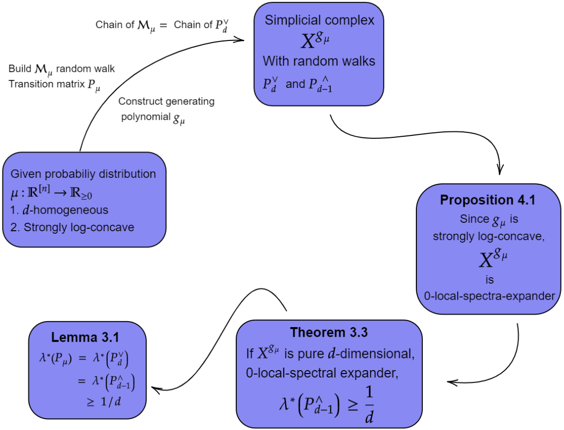

Figure 7: Roadmap

References

Anari, N., Liu, K., Gharan, S. O., Vinzant, C. (2019, June). *Log-concave polynomials II: High-dimensional walks and an FPRAS for counting bases of a matroid*. In Proceedings of the 51st Annual ACM SIGACT Symposium on Theory of Computing (pp. 1-16).

Mark Jerrum, Leslie Valiant, and Vijay Vazirani. "Random Generation of Combinatorial Structures from a Uniform Distribution". In: Theoretical Computer Science 43 (1986), pp. 169–188.

Nima Anari, Shayan Oveis Gharan, and Cynthia Vinzant. "Log-concave polynomials, entropy, and a deterministic approximation algorithm for counting bases of matroids". In: *FOCS*. to appear. 2018.

Alexander Barvinok and Alex Samorodnitsky. "Random weighting, asymptotic counting, and inverse isoperimetry". In: Israel Journal of Mathematics 158.1 (Mar. 2007), pp. 159–191. issn: 1565-8511.